

# proj\_walmart

April 9, 2024

## Importing Libraries

```
[109]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

## Loading the Dataset .csv file

```
[110]: df = pd.read_csv('Walmart_Store_sales.csv')
```

```
[111]: df
```

```
[111]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	\
0	1	05-02-2010	1643690.90	0	42.31	2.572	
1	1	12-02-2010	1641957.44	1	38.51	2.548	
2	1	19-02-2010	1611968.17	0	39.93	2.514	
3	1	26-02-2010	1409727.59	0	46.63	2.561	
4	1	05-03-2010	1554806.68	0	46.50	2.625	
...	...	...	...	...	...	...	
6430	45	28-09-2012	713173.95	0	64.88	3.997	
6431	45	05-10-2012	733455.07	0	64.89	3.985	
6432	45	12-10-2012	734464.36	0	54.47	4.000	
6433	45	19-10-2012	718125.53	0	56.47	3.969	
6434	45	26-10-2012	760281.43	0	58.85	3.882	

	CPI	Unemployment
0	211.096358	8.106
1	211.242170	8.106
2	211.289143	8.106
3	211.319643	8.106
4	211.350143	8.106
...	...	...
6430	192.013558	8.684
6431	192.170412	8.667
6432	192.327265	8.667

```
6433  192.330854      8.667
6434  192.308899      8.667
```

```
[6435 rows x 8 columns]
```

```
[112]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Store            6435 non-null   int64
1   Date             6435 non-null   object
2   Weekly_Sales     6435 non-null   float64
3   Holiday_Flag     6435 non-null   int64
4   Temperature      6435 non-null   float64
5   Fuel_Price       6435 non-null   float64
6   CPI              6435 non-null   float64
7   Unemployment     6435 non-null   float64
dtypes: float64(5), int64(2), object(1)
memory usage: 402.3+ KB
```

```
[113]: df.shape
```

```
[113]: (6435, 8)
```

- The number of Rows and Columns are 6435 and 8

```
[114]: df.describe()
```

```
[114]:
```

	Store	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price \
count	6435.000000	6.435000e+03	6435.000000	6435.000000	6435.000000
mean	23.000000	1.046965e+06	0.069930	60.663782	3.358607
std	12.988182	5.643666e+05	0.255049	18.444933	0.459020
min	1.000000	2.099862e+05	0.000000	-2.060000	2.472000
25%	12.000000	5.533501e+05	0.000000	47.460000	2.933000
50%	23.000000	9.607460e+05	0.000000	62.670000	3.445000
75%	34.000000	1.420159e+06	0.000000	74.940000	3.735000
max	45.000000	3.818686e+06	1.000000	100.140000	4.468000

	CPI	Unemployment
count	6435.000000	6435.000000
mean	171.578394	7.999151
std	39.356712	1.875885
min	126.064000	3.879000
25%	131.735000	6.891000
50%	182.616521	7.874000

75%	212.743293	8.622000
max	227.232807	14.313000

```
[115]: df.isna().sum()
```

```
[115]: Store          0
      Date          0
      Weekly_Sales  0
      Holiday_Flag  0
      Temperature  0
      Fuel_Price    0
      CPI           0
      Unemployment  0
      dtype: int64
```

- There is no missing values present in the dataset

```
[116]: df.dtypes
```

```
[116]: Store          int64
      Date          object
      Weekly_Sales  float64
      Holiday_Flag  int64
      Temperature  float64
      Fuel_Price    float64
      CPI           float64
      Unemployment  float64
      dtype: object
```

- As checked, the Date variable has object data-type

```
[117]: # Converting the Date data-type to datetime
```

```
[118]: from datetime import datetime

      df['Date'] = pd.to_datetime(df['Date'])
```

```
[119]: df.dtypes
```

```
[119]: Store          int64
      Date          datetime64[ns]
      Weekly_Sales  float64
      Holiday_Flag  int64
      Temperature  float64
      Fuel_Price    float64
      CPI           float64
      Unemployment  float64
      dtype: object
```

```
[120]: df
```

```
[120]:      Store      Date  Weekly_Sales  Holiday_Flag  Temperature  Fuel_Price  \
0         1 2010-05-02   1643690.90             0         42.31         2.572
1         1 2010-12-02   1641957.44             1         38.51         2.548
2         1 2010-02-19   1611968.17             0         39.93         2.514
3         1 2010-02-26   1409727.59             0         46.63         2.561
4         1 2010-05-03   1554806.68             0         46.50         2.625
...     ...      ...      ...      ...      ...      ...
6430      45 2012-09-28    713173.95             0         64.88         3.997
6431      45 2012-05-10    733455.07             0         64.89         3.985
6432      45 2012-12-10    734464.36             0         54.47         4.000
6433      45 2012-10-19    718125.53             0         56.47         3.969
6434      45 2012-10-26    760281.43             0         58.85         3.882

      CPI  Unemployment
0    211.096358         8.106
1    211.242170         8.106
2    211.289143         8.106
3    211.319643         8.106
4    211.350143         8.106
...     ...      ...
6430  192.013558         8.684
6431  192.170412         8.667
6432  192.327265         8.667
6433  192.330854         8.667
6434  192.308899         8.667

[6435 rows x 8 columns]
```

## 0.1 Basic Statistics tasks

### 1. Which store has maximum sales

```
[121]: total_sales=df.groupby('Store')['Weekly_Sales'].sum('Weekly_Sales').round()

      # Grouping the Store number and its Weekly Sales to finding the sum of all
      ↪Weekly Sales of each Store
```

```
[122]: pd.DataFrame(total_sales.sort_values().tail(1))

      # Sort the values and find the maximum Weekly Sales of the Store
```

```
[122]:      Weekly_Sales
Store
20      301397792.0
```

- Here, the Store 20 has the maximum sales

**2. Which store has maximum standard deviation i.e., the sales vary a lot. Also, find out the coefficient of mean to standard deviation**

```
[123]: df_std = df.groupby('Store')['Weekly_Sales'].std()
```

```
[124]: pd.DataFrame(df_std.sort_values().tail(1))
```

```
[124]:      Weekly_Sales
Store
14      317569.949476
```

- The Store 14 has the maximum standard deviation

```
[125]: store_max = df[df['Store']==14].Weekly_Sales

mean_to_std = store_max.std() / store_max.mean()*100
```

```
[126]: mean_to_std
```

```
[126]: 15.713673600948338
```

- The co-efficient of mean to std deviation of store 14 is 15.713%

**3. Which store/s has good quarterly growth rate in Q3'2012**

```
[127]: df.head()
```

```
[127]:      Store      Date  Weekly_Sales  Holiday_Flag  Temperature  Fuel_Price  \
0         1  2010-05-02    1643690.90             0         42.31         2.572
1         1  2010-12-02    1641957.44             1         38.51         2.548
2         1  2010-02-19    1611968.17             0         39.93         2.514
3         1  2010-02-26    1409727.59             0         46.63         2.561
4         1  2010-05-03    1554806.68             0         46.50         2.625
```

```
      CPI  Unemployment
0  211.096358         8.106
1  211.242170         8.106
2  211.289143         8.106
3  211.319643         8.106
4  211.350143         8.106
```

```
[128]: q2 = df[(df['Date']>='2012-04-01')&(df['Date']<='2012-06-30')].
        ↳groupby('Store')['Weekly_Sales'].sum().round()
```

```
[129]: q2.head()
```

```
[129]: Store
1      21036966.0
2      25085124.0
```

```

3      5562668.0
4      28384185.0
5      4427262.0
Name: Weekly_Sales, dtype: float64

```

```
[130]: q3 = df[(df['Date']>='2012-07-01')&(df['Date']<='2012-09-30')].
        ↳groupby('Store')['Weekly_Sales'].sum().round()
```

```
[131]: q3.head()
```

```
[131]: Store
1      18633210.0
2      22396868.0
3       4966496.0
4      25652119.0
5       3880622.0
Name: Weekly_Sales, dtype: float64

```

```
[132]: df_q3_2012 = pd.DataFrame({'Q2 Sales':q2,'Q3 Sales':q3,'Difference':(q3 - q2),
                                  'Growth Rate':(q3-q2)/q2})

df_q3_2012.sort_values(by = 'Growth Rate', ascending=False).head(1)

# Formula of Growth Rate is (q3 - q2)/q2
# I have taken the top store which has good Growth Rate among all

```

```
[132]:      Q2 Sales    Q3 Sales  Difference  Growth Rate
Store
16      6626133.0  6441311.0   -184822.0   -0.027893

```

- Here, the store 16 has good quarterly growth rate in Q3'2012 which has a Growth Rate of -0.027

**4. Some holidays have a negative impact on sales. Find out holidays which have higher sales than the mean sales in non-holiday season for all stores together**

```
[133]: df
```

```
[133]:      Store      Date  Weekly_Sales  Holiday_Flag  Temperature  Fuel_Price  \
0         1  2010-05-02    1643690.90             0         42.31         2.572
1         1  2010-12-02    1641957.44             1         38.51         2.548
2         1  2010-02-19    1611968.17             0         39.93         2.514
3         1  2010-02-26    1409727.59             0         46.63         2.561
4         1  2010-05-03    1554806.68             0         46.50         2.625
...     ...      ...
6430    45  2012-09-28     713173.95             0         64.88         3.997
6431    45  2012-05-10     733455.07             0         64.89         3.985
6432    45  2012-12-10     734464.36             0         54.47         4.000

```

6433	45	2012-10-19	718125.53	0	56.47	3.969
6434	45	2012-10-26	760281.43	0	58.85	3.882

	CPI	Unemployment
0	211.096358	8.106
1	211.242170	8.106
2	211.289143	8.106
3	211.319643	8.106
4	211.350143	8.106
...	...	...
6430	192.013558	8.684
6431	192.170412	8.667
6432	192.327265	8.667
6433	192.330854	8.667
6434	192.308899	8.667

[6435 rows x 8 columns]

```
[134]: Super_Bowl=['12-02-2010','11-02-2011','10-02-2012','08-02-2013']
Labour_Day = ['2010-09-10','2011-09-09','2012-09-07','2013-09-06']
Thanksgiving = ['2010-11-26','2011-11-25','2012-11-23','2013-11-29']
Christmas = ['2010-12-31','2011-12-30','2012-12-28','2013-12-27']
```

```
# Created a set from the holiday events description
```

```
[135]: Super_Bowl_Sales = round(df[df['Date'].isin(Super_Bowl)]['Weekly_Sales'].
↳mean(),2)
Labour_Day_Sales = round(df[df['Date'].isin(Labour_Day)]['Weekly_Sales'].
↳mean(),2)
Thanksgiving_Sales = round(df[df['Date'].isin(Thanksgiving)]['Weekly_Sales'].
↳mean(),2)
Christmas_Sales = round(df[df['Date'].isin(Christmas)]['Weekly_Sales'].mean(),2)
non_holiday_sales = round(df[df['Holiday_Flag']==0]['Weekly_Sales'].mean(),2)
```

```
[136]: print('Super bowl: ' + str(Super_Bowl_Sales))
print('Labour Day: ' + str(Labour_Day_Sales))
print('Thanksgiving: ' + str(Thanksgiving_Sales))
print('Christmas: ' + str(Christmas_Sales))
print('Non-Holiday sales: ' +str(non_holiday_sales))
```

```
Super bowl: 1079127.99
Labour Day: 1039182.83
Thanksgiving: 1471273.43
Christmas: 960833.11
Non-Holiday sales: 1041256.38
```

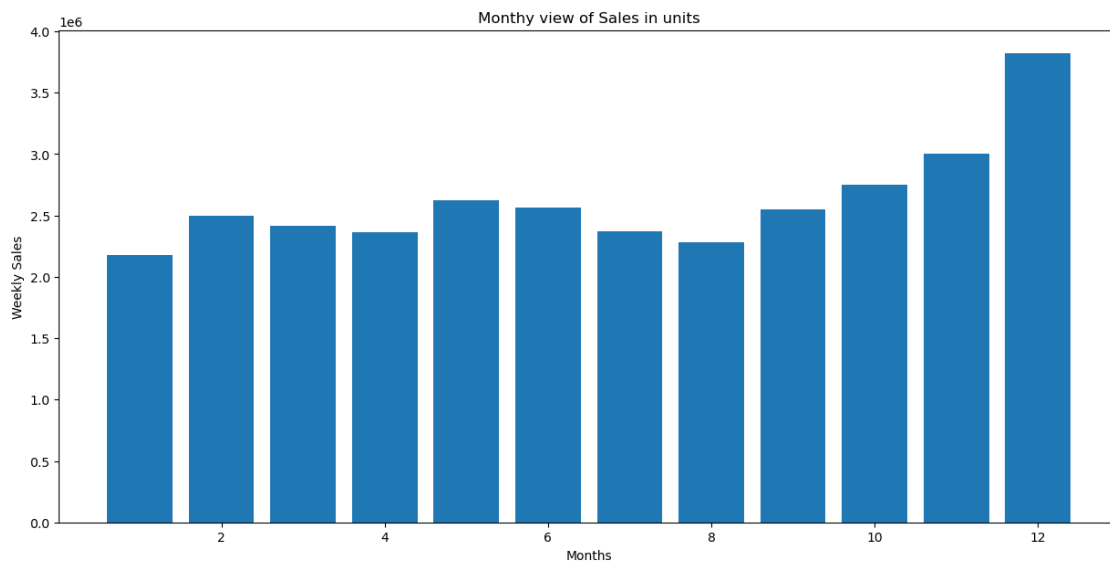
- Here, Thanksgiving has the highest sales than the mean sales in non-holiday season for all stores

## 5. Provide a monthly and semester view of sales in units and give insights

```
[137]: # Monthly view of Sales in units
```

```
[138]: df['year'] = pd.DatetimeIndex(df['Date']).year  
df['month'] = pd.DatetimeIndex(df['Date']).month  
df['day'] = pd.DatetimeIndex(df['Date']).day
```

```
[139]: plt.figure(figsize=(15,7))  
  
plt.bar(df['month'],df['Weekly_Sales'])  
plt.xlabel('Months')  
plt.ylabel('Weekly Sales')  
  
plt.title('Monthly view of Sales in units')  
plt.show()
```



- Here, December has the highest sales

```
[140]: # Semester view of Sales in units
```

```
[141]: df['semester'] = np.where(df['month'] < 7, 1, 2)  
df['semester']
```

```
[141]: 0      1  
      1      2  
      2      1  
      3      1  
      4      1
```



```


..
6430    2
6431    1
6432    2
6433    2
6434    2
Name: semester, Length: 6435, dtype: int32

```

```

[142]: semester=sns.barplot(x='semester',y='Weekly_Sales',data=df)
plt.show()

```



- Here, Semester 2 has the highest sales in unit

## 0.2 Statistical Model

```

[143]: X = df[['Store','Fuel_Price','CPI','Unemployment','day','month','year']]
Y = df['Weekly_Sales']

```

```

[144]: # Linear Regression model

```

```
[145]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import sklearn
```

```
[146]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, train_size=0.8,
↳random_state=42)
```

```
[147]: linear_reg = LinearRegression()
```

```
[148]: linear_reg.fit(X_train, Y_train)
```

```
[148]: LinearRegression()
```

```
[151]: Y_pred = linear_reg.predict(X_test)
```

```
[152]: print('R2 score: ' +str(r2_score(Y_test, Y_pred)))
print('Linear model accuracy: ' +str(linear_reg.score(X_train, Y_train)))
print('Mean Squared Error: ' +str(mean_squared_error(Y_test, Y_pred)))
```

R2 score: 0.14894500845355385

Linear model accuracy: 0.14372803259754718

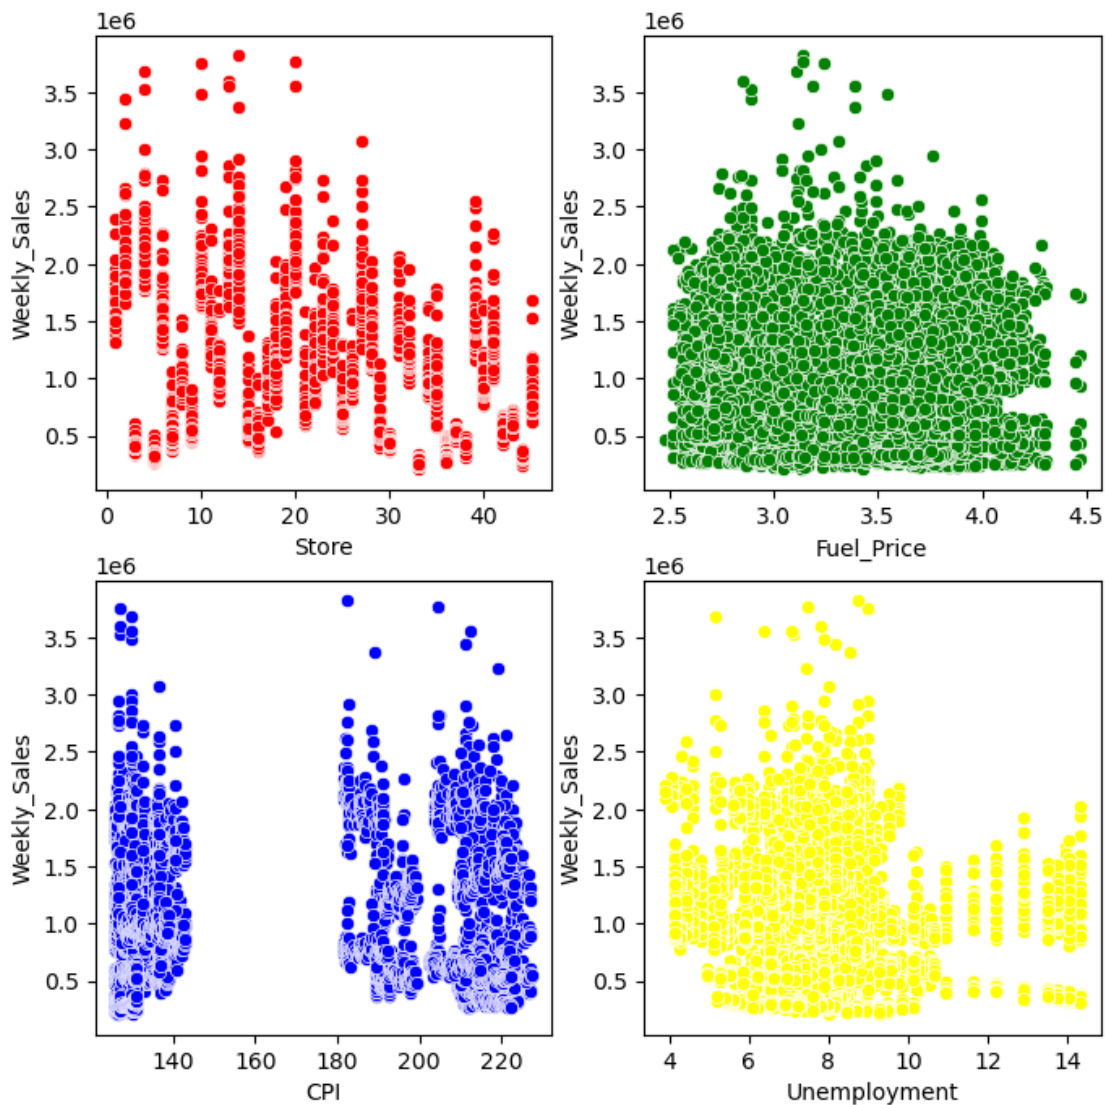
Mean Squared Error: 274171250281.01086

```
[153]: # Plot a scatterplot graph
```

```
[154]: figs,axes = plt.subplots(2,2, figsize = (8,8))

sns.scatterplot(data=df, x='Store', y='Weekly_Sales', color='red', ax=axes[0,0])
sns.scatterplot(data=df, x='Fuel_Price', y='Weekly_Sales', color='green',
↳ax=axes[0,1])
sns.scatterplot(data=df, x='CPI', y='Weekly_Sales', color='blue', ax=axes[1,0])
sns.scatterplot(data=df, x='Unemployment', y='Weekly_Sales', color='yellow',
↳ax=axes[1,1])

plt.show()
```



- Linear Regression is not good fit for the model

```
[155]: # Random Forest Model
```

```
[156]: df.corr()
```

```
[156]:
```

	Store	Weekly_Sales	Holiday_Flag	Temperature \
Store	1.000000e+00	-0.335332	-4.386841e-16	-0.022659
Weekly_Sales	-3.353320e-01	1.000000	3.689097e-02	-0.063810
Holiday_Flag	-4.386841e-16	0.036891	1.000000e+00	-0.155091
Temperature	-2.265908e-02	-0.063810	-1.550913e-01	1.000000
Fuel_Price	6.002295e-02	0.009464	-7.834652e-02	0.144982
CPI	-2.094919e-01	-0.072634	-2.162091e-03	0.176888

Unemployment	2.235313e-01	-0.106176	1.096028e-02	0.101158
year	3.474318e-12	-0.018378	-5.678257e-02	0.064269
month	6.289676e-16	0.067535	3.322341e-01	0.066440
day	-1.070464e-15	-0.014873	-3.603594e-02	0.089019
semester	5.868729e-16	0.035353	2.761285e-01	0.130314

	Fuel_Price	CPI	Unemployment	year	month \
Store	0.060023	-0.209492	0.223531	3.474318e-12	6.289676e-16
Weekly_Sales	0.009464	-0.072634	-0.106176	-1.837754e-02	6.753523e-02
Holiday_Flag	-0.078347	-0.002162	0.010960	-5.678257e-02	3.322341e-01
Temperature	0.144982	0.176888	0.101158	6.426923e-02	6.643970e-02
Fuel_Price	1.000000	-0.170642	-0.034684	7.794703e-01	-5.283174e-02
CPI	-0.170642	1.000000	-0.302020	7.479573e-02	1.478843e-03
Unemployment	-0.034684	-0.302020	1.000000	-2.418135e-01	-2.061552e-03
year	0.779470	0.074796	-0.241813	1.000000e+00	-1.390145e-01
month	-0.052832	0.001479	-0.002062	-1.390145e-01	1.000000e+00
day	0.032532	0.003966	-0.008167	-1.277942e-02	5.959249e-03
semester	-0.061948	0.002291	-0.009498	-9.683011e-02	8.642165e-01

	day	semester
Store	-1.070464e-15	5.868729e-16
Weekly_Sales	-1.487292e-02	3.535312e-02
Holiday_Flag	-3.603594e-02	2.761285e-01
Temperature	8.901925e-02	1.303141e-01
Fuel_Price	3.253169e-02	-6.194830e-02
CPI	3.965821e-03	2.291341e-03
Unemployment	-8.166853e-03	-9.497961e-03
year	-1.277942e-02	-9.683011e-02
month	5.959249e-03	8.642165e-01
day	1.000000e+00	3.487338e-02
semester	3.487338e-02	1.000000e+00

```
[157]: from sklearn.ensemble import RandomForestRegressor
```

```
[175]: rf = RandomForestRegressor(n_estimators = 500, max_depth=15, n_jobs=5)
       rf.fit(X_train,Y_train)
```

```
[175]: RandomForestRegressor(max_depth=15, n_estimators=500, n_jobs=5)
```

```
[176]: Y_pred=rf.predict(X_test)
```

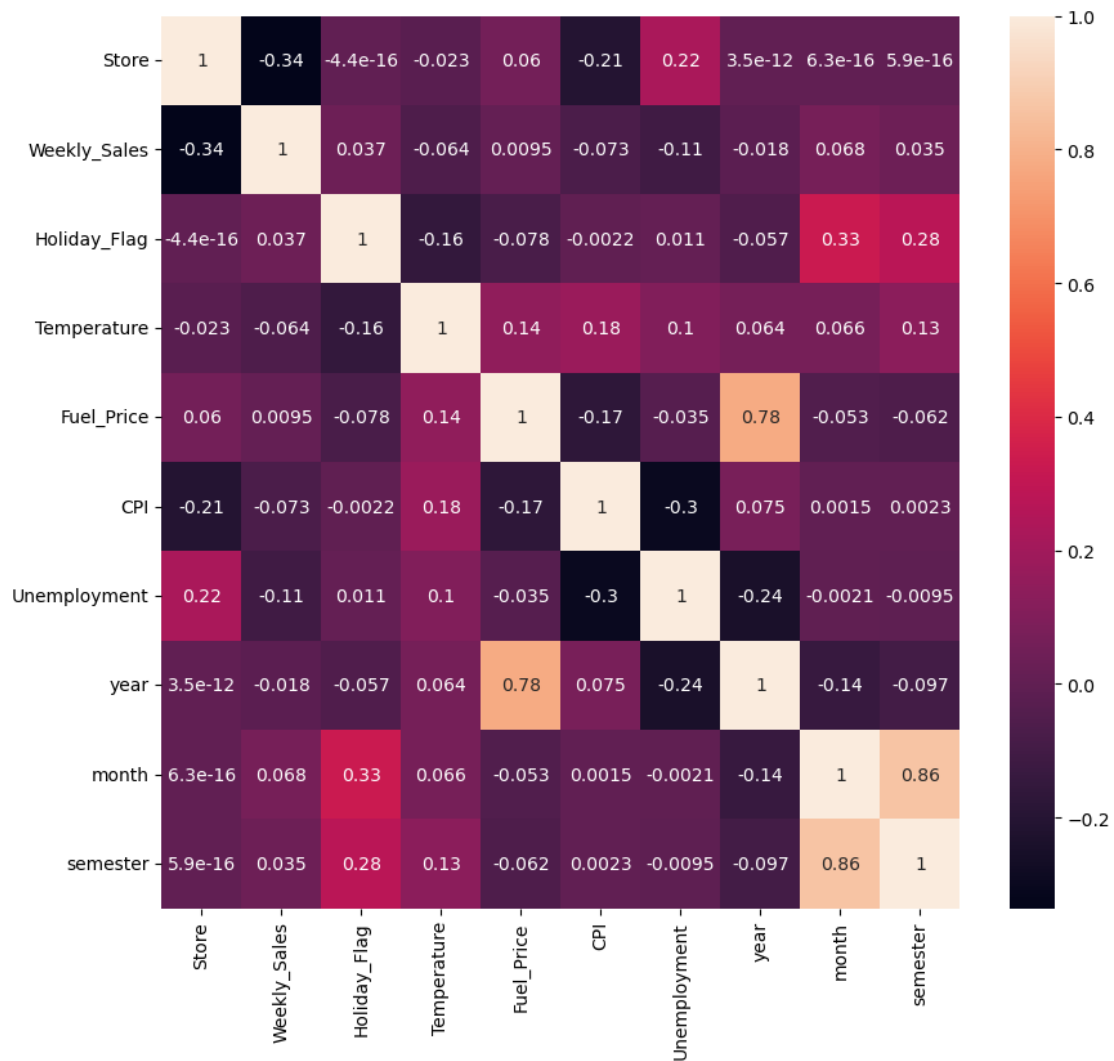
```
[177]: print('R2 score: ' +str(r2_score(Y_test, Y_pred)))
       print('Root Mean Squared Error:', np.sqrt(mean_squared_error(Y_test, Y_pred)))
       print('Mean Squared Error: ' +str(mean_squared_error(Y_test, Y_pred)))
```

R2 score: 0.9524234944111671


Root Mean Squared Error: 123802.21502366115

Mean Squared Error: 15326988444.76483

```
[178]: plt.figure(figsize=(10,9))
sns.heatmap(df.corr(),annot=True)
plt.show()
```



```
[179]: plt.scatter(Y_test, Y_pred)
plt.xlabel('Y_test')
plt.ylabel('Prediction')
plt.show()
```



- For Random Forest Regressor R2 score: R2 score: 0.9524234944111671 Root Mean Squared Error: 123802.21502366115 Mean Squared Error: 15326988444.76483
- The Random Forest Regressor model would be the best fit for the outcome

## 2. Change dates into days by creating new variable.

```
[164]: df['day'] = pd.to_datetime(df['Date']).dt.day_name()
df.head()
```

```
[164]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	\
0	1	2010-05-02	1643690.90	0	42.31	2.572	
1	1	2010-12-02	1641957.44	1	38.51	2.548	
2	1	2010-02-19	1611968.17	0	39.93	2.514	
3	1	2010-02-26	1409727.59	0	46.63	2.561	
4	1	2010-05-03	1554806.68	0	46.50	2.625	

	CPI	Unemployment	year	month	day	semester
0	211.096358	8.106	2010	5	Sunday	1
1	211.242170	8.106	2010	12	Thursday	2
2	211.289143	8.106	2010	2	Friday	1

3	211.319643	8.106	2010	2	Friday	1
4	211.350143	8.106	2010	5	Monday	1

[ ]:

[ ]: